

*Let's Go Asynchronous*



# Tomáš Sedláček

[mail@kedlas.cz](mailto:mail@kedlas.cz)  
<https://www.linkedin.com/in/tomasedlacek/>

Q&A



Monolithic app

~~Monolithic app~~

Synchronously communicating microservices

~~Monolithic app~~

~~Synchronously communicating microservices~~

*Asynchronously communicating microservices*

*AWS Managed RabbitMQ*

~~Monolithic app~~

~~Synchronously communicating microservices~~

**Asynchronously communicating microservices**

~~Managed RabbitMQ~~

**PGQ**



*Meet Orafo*

# Order process

- get and validate items
- get and validate customer details
- apply coupons
- create the order db record
- generate invoice
- process payment
- order shipping
- synchronize with CRM/ERP systems
- ...

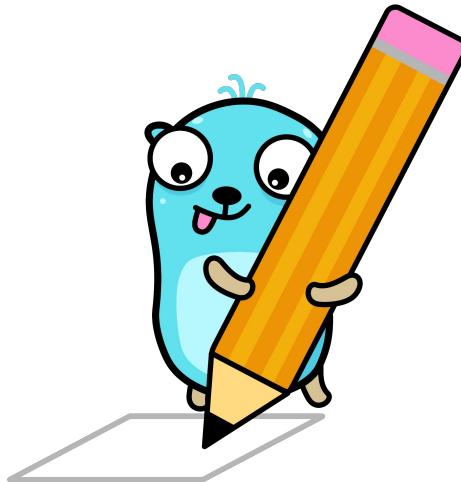
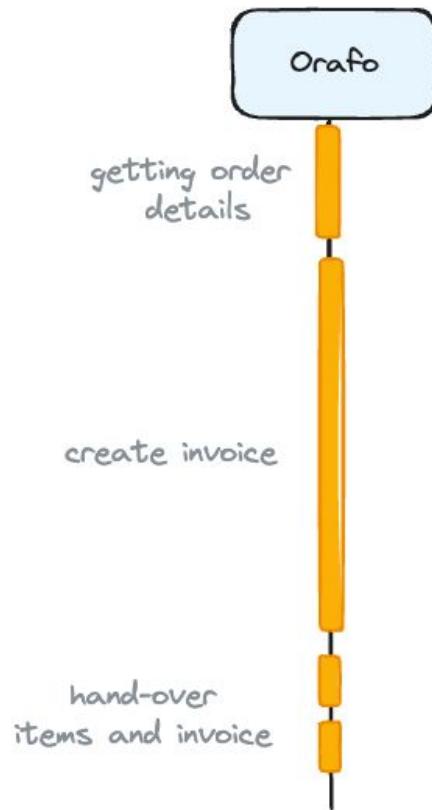


# *Simplified order process*

1. *get order details (validate items & customer)*
2. *generate invoice*
3. *hand-over*



Synchronous  
processing



```
type orderDetails struct{}

type invoice struct{}

func main() {
    http.HandleFunc("/order", handleOrder)
    fmt.Println("Server listening on port 8090")
    err := http.ListenAndServe(":8090", nil)
    if err != nil {
        panic(err)
    }
}

func handleOrder(w http.ResponseWriter, _ *http.Request) {
    details := getOrderDetails()
    inv := createInvoice(details)
    handover(details, inv)
    _, _ = fmt.Fprint(w, "Order processed")
}

func getOrderDetails() orderDetails {
    time.Sleep(100 * time.Millisecond)
    return orderDetails{}
}

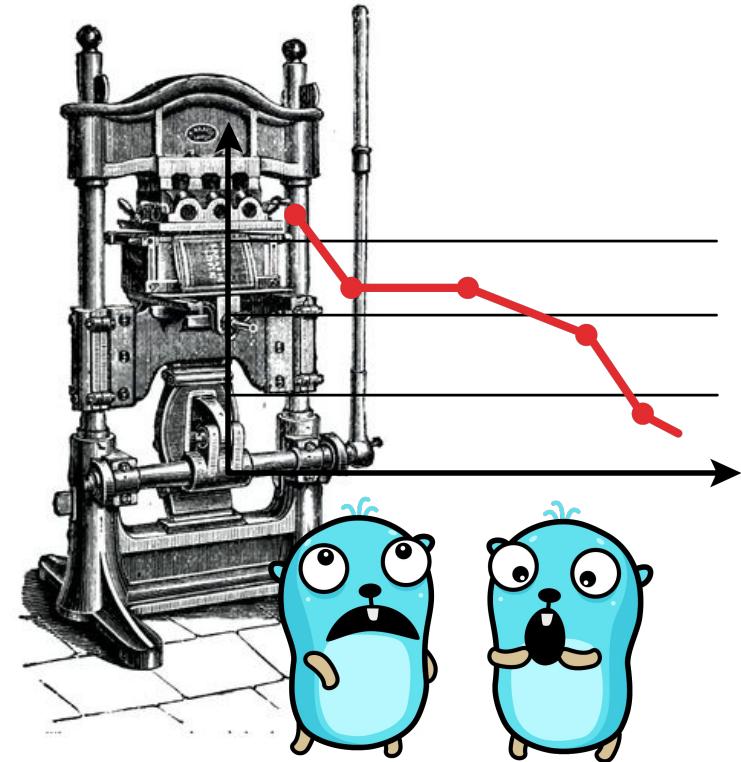
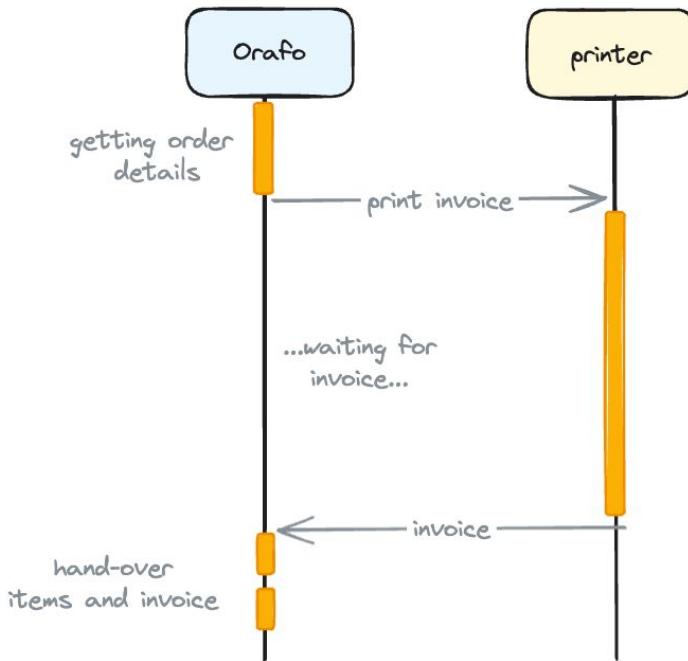
func createInvoice(_ orderDetails) invoice {
    time.Sleep(5 * time.Second)
    return invoice{}
}

func handover(_ orderDetails, _ invoice) {
    time.Sleep(200 * time.Millisecond)
}
```

sync/http



Synchronous  
processing



```

func main() {
    http.HandleFunc("/order", handleOrder)
    fmt.Println("Server listening on port 8090")
    err := http.ListenAndServe(":8090", nil)
    if err != nil { panic(err) }
}

func handleOrder(w http.ResponseWriter, _ *http.Request) {
    details := getOrderDetails()
    err, inv := createInvoice(details)
    if err != nil {
        _, _ = fmt.Fprint(w, "Order process failed: failed to create invoice")
        return
    }
    handover(details, inv)
    _, _ = fmt.Fprint(w, "Order processed")
}

func createInvoice(details orderDetails) (error, invoice) {
    reqBodyBytes := new(bytes.Buffer)
    _ = json.NewEncoder(reqBodyBytes).Encode(details)

    resp, err := http.Post(
        url: "http://127.0.0.1:8091/print",
        contentType: "application/json",
        reqBodyBytes,
    )
    if err != nil {
        return err, invoice{}
    }

    return nil, parsePrinterResponse(resp)
}

```

```

func main() {
    http.HandleFunc("/print", handlePrint)
    fmt.Println("Printer listening on port 8091")
    _ = http.ListenAndServe(":8091", nil)
}

func handlePrint(w http.ResponseWriter, req *http.Request) {
    details := parseRequest(req)
    inv := printInvoice(details)

    _, _ = fmt.Fprint(w, inv)
}

func parseRequest(_ *http.Request) orderDetails { return orderDetails{} }
func printInvoice(_ orderDetails) invoice {
    return invoice{}
}

```

# sync/http-printer



# *synchronous code*

- simple to write, read and debug
- ordered, goes function by function
- total time is the sum of each function times
- usually highly I/O dependant
- no concurrency, unless you use goroutines, async/await or similar

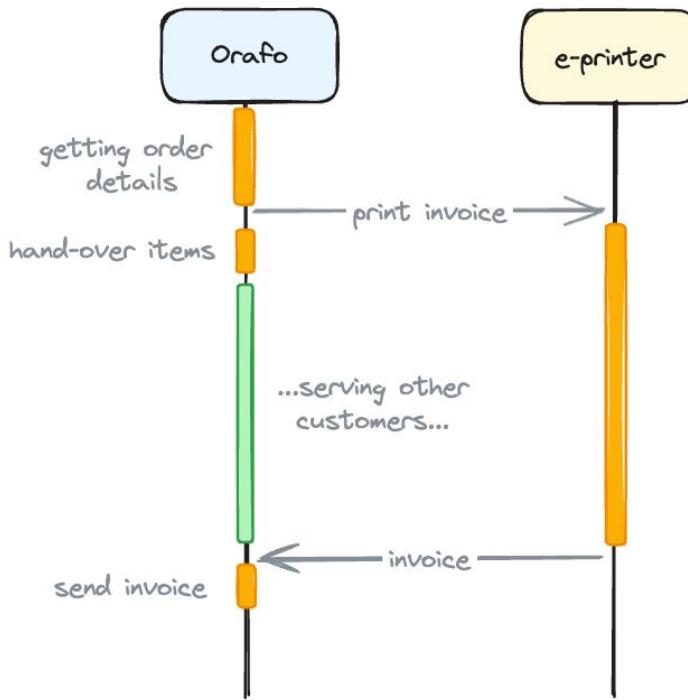
# *synchronous code*

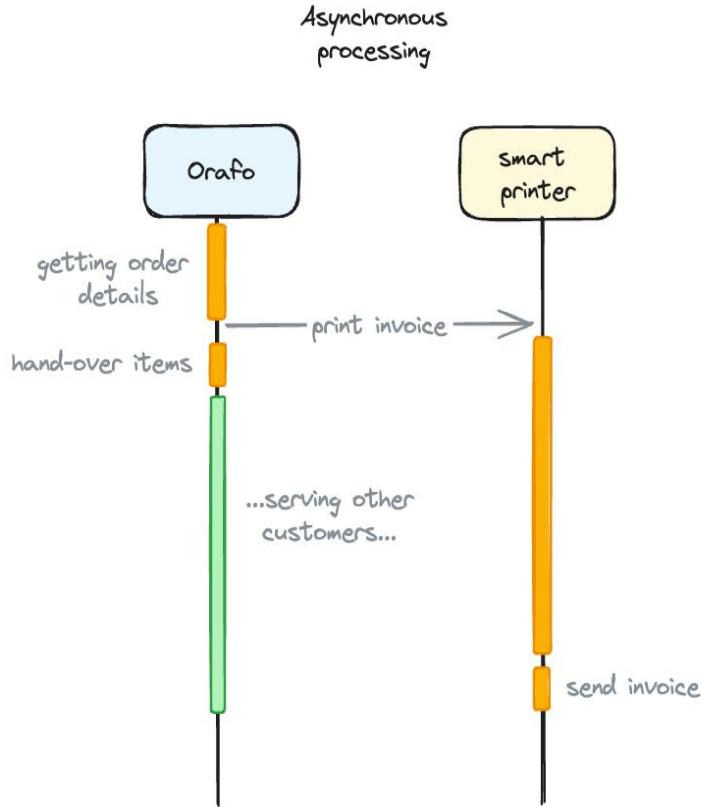
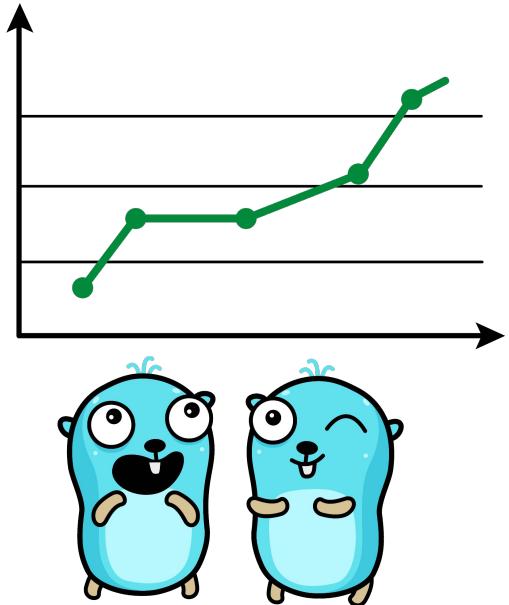
- simple to write, read and debug
- ordered, goes function by function
- total time is the sum of each function times
- highly I/O dependant
- no concurrency, unless you use goroutines, async/await or similar
- on failures you must rollback, retry or ignore error
  - printer is offline
  - printer crash
  - printer out of ink, printer busy, printer timeout, ...

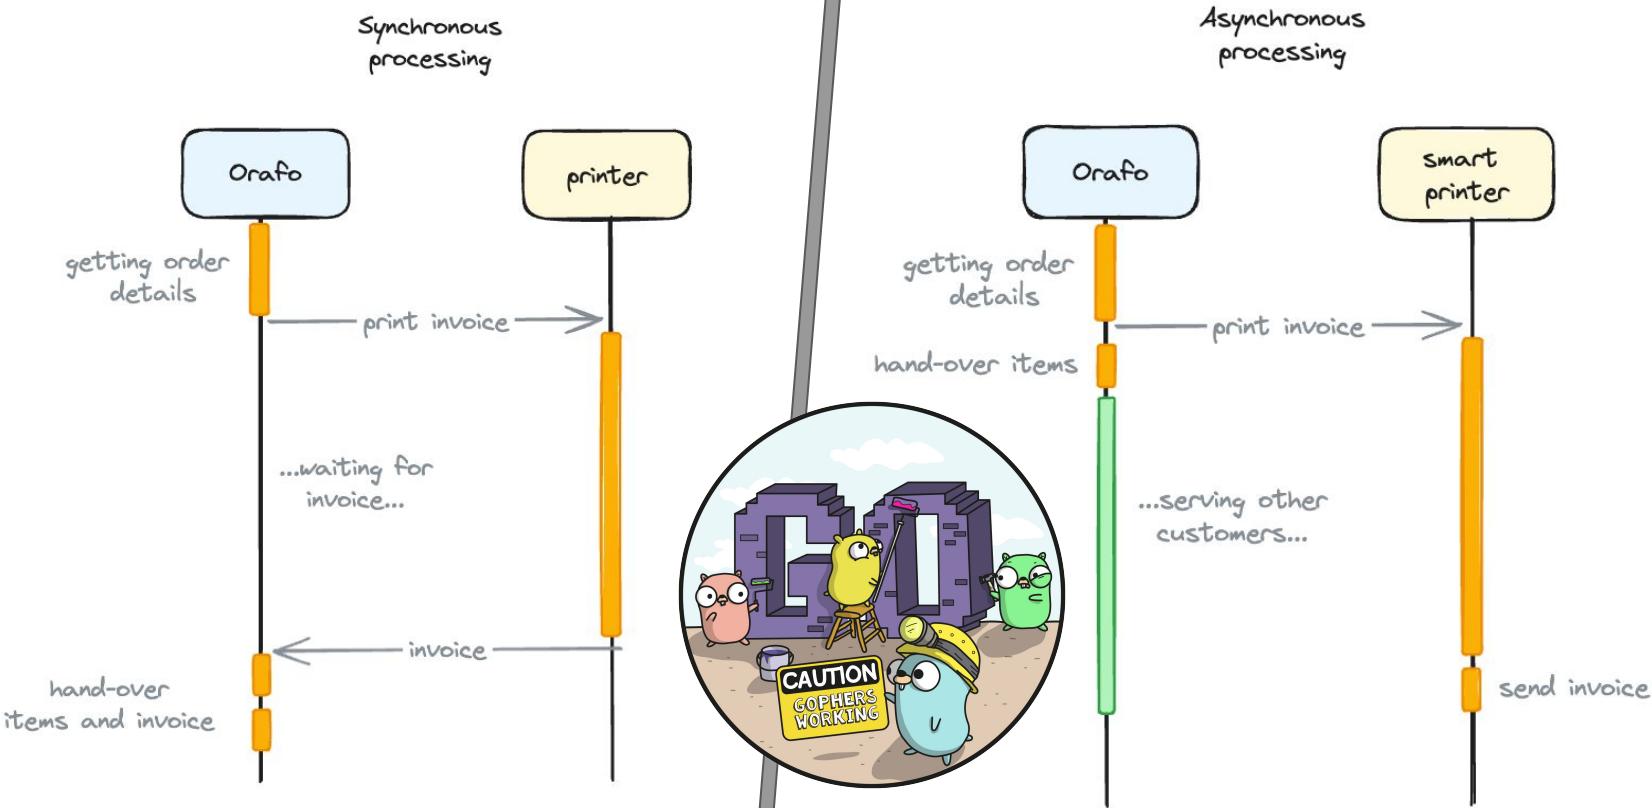
*Let's Go Asynchronous*



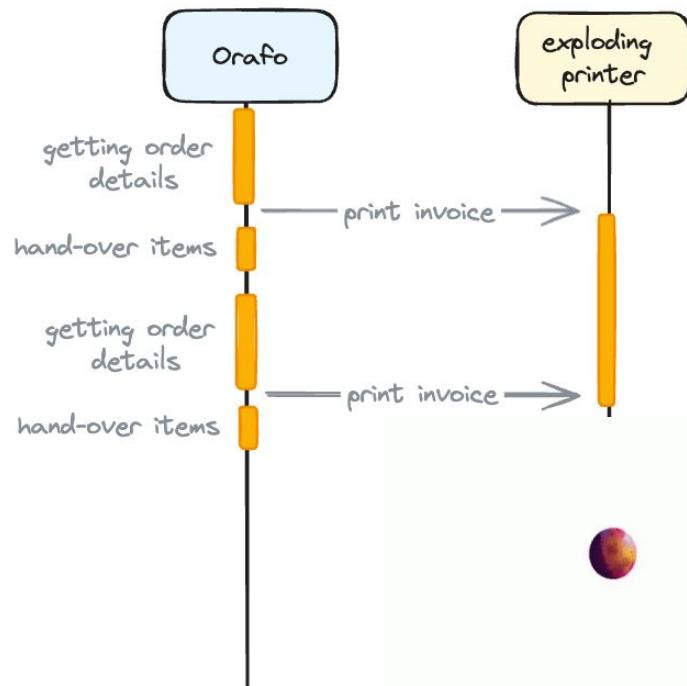
Asynchronous  
processing



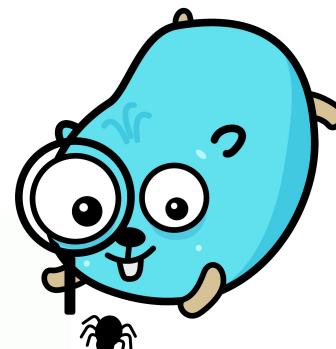




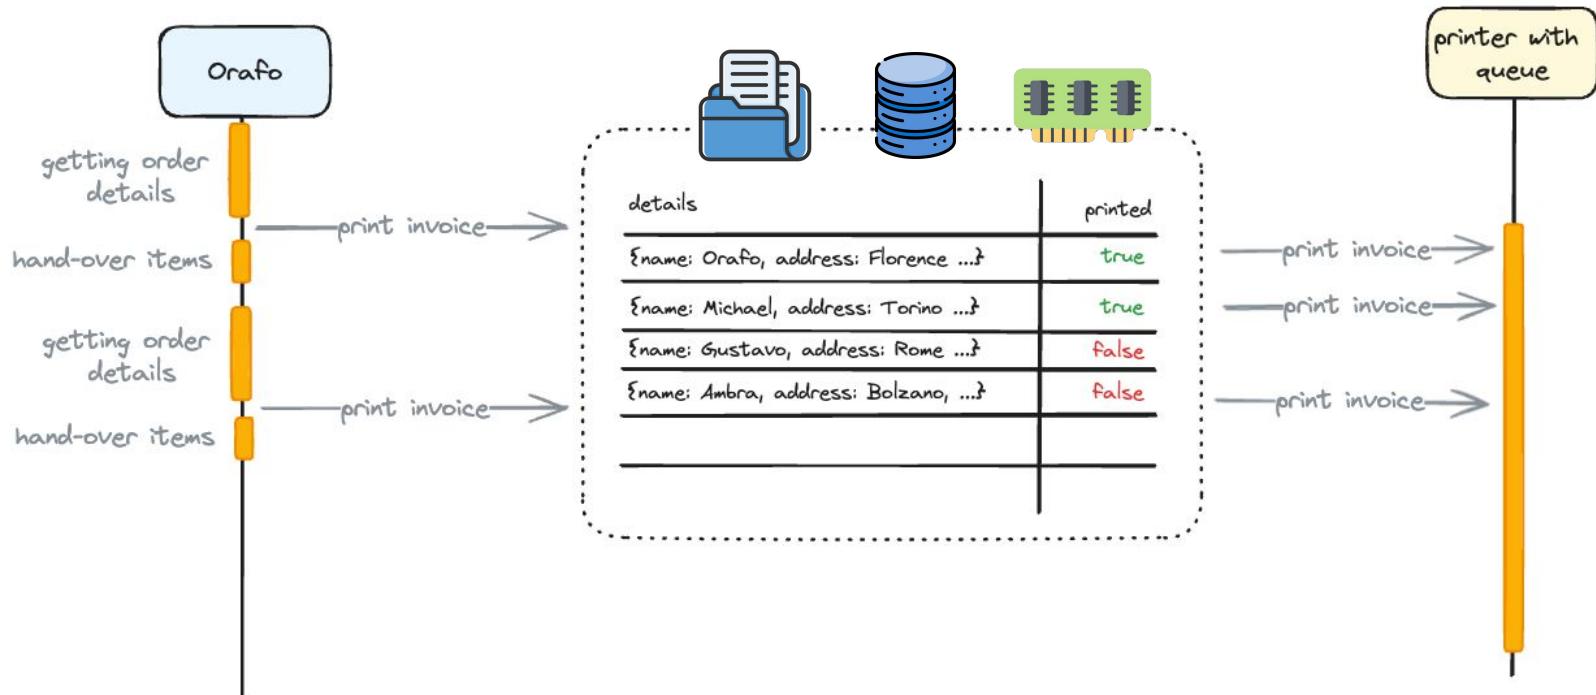
Resources limits



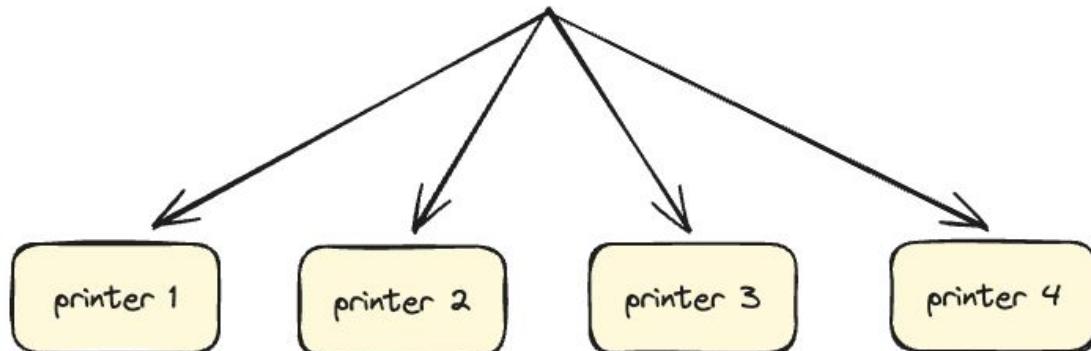
# async/http/printer



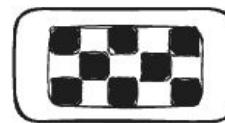
DEMO



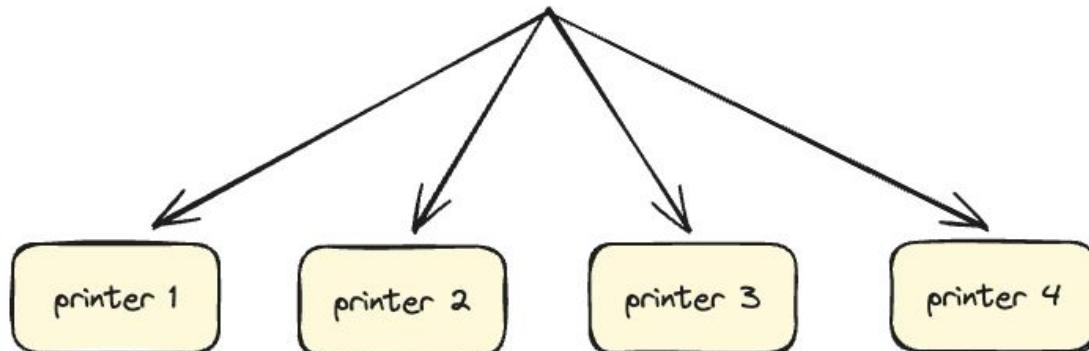
details	printed
{ename: Orafo, address: Florence ...}	true
{ename: Michael, address: Torino ...}	true
{ename: Gustavo, address: Rome ...}	false
{ename: Ambra, address: Bolzano, ...}	false
...	...
$n > \boxed{XXL}$	false



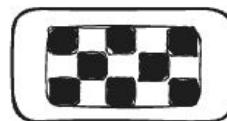
details	printed
{ename: Orafo, address: Florence ...}	true
{ename: Michael, address: Torino ...}	true
{ename: Gustavo, address: Rome ...}	false
{ename: Ambra, address: Bolzano, ...}	false
...	...
$n > \boxed{\text{XXL}}$	false



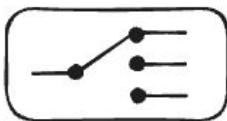
race conditions



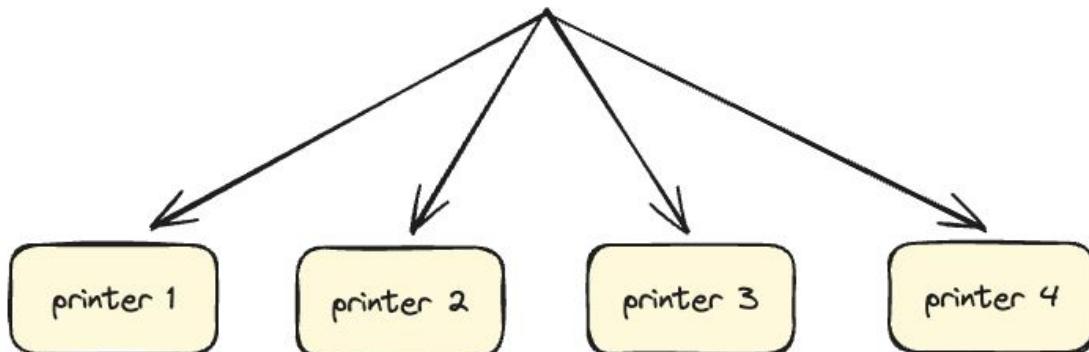
details	printed
{ename: Orafo, address: Florence ...}	true
{ename: Michael, address: Torino ...}	true
{ename: Gustavo, address: Rome ...}	false
{ename: Ambra, address: Bolzano, ...}	false
...	...
$n > \boxed{\text{XXL}}$	false



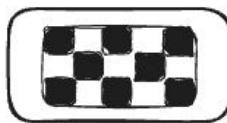
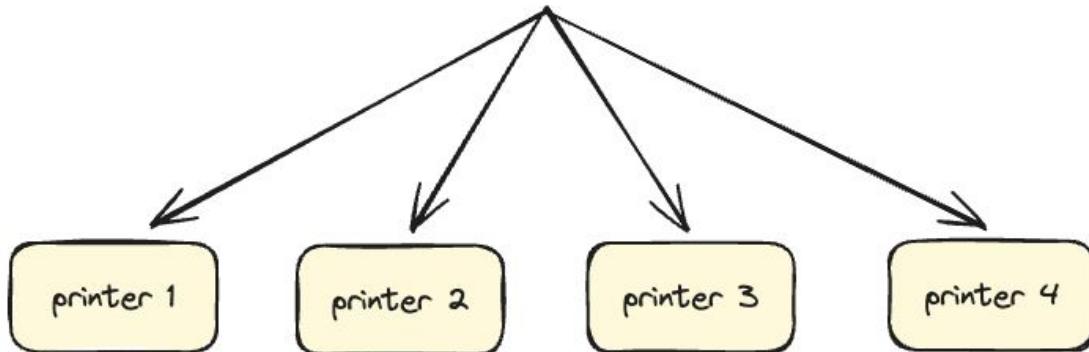
race conditions



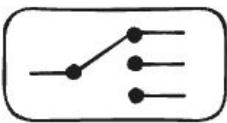
routing



details	printed
{ename: Orafo, address: Florence ...}	true
{ename: Michael, address: Torino ...}	true
{ename: Gustavo, address: Rome ...}	false
{ename: Ambra, address: Bolzano, ...}	false
...	...
$n > \boxed{XXL}$	false



race conditions

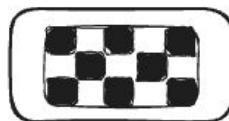
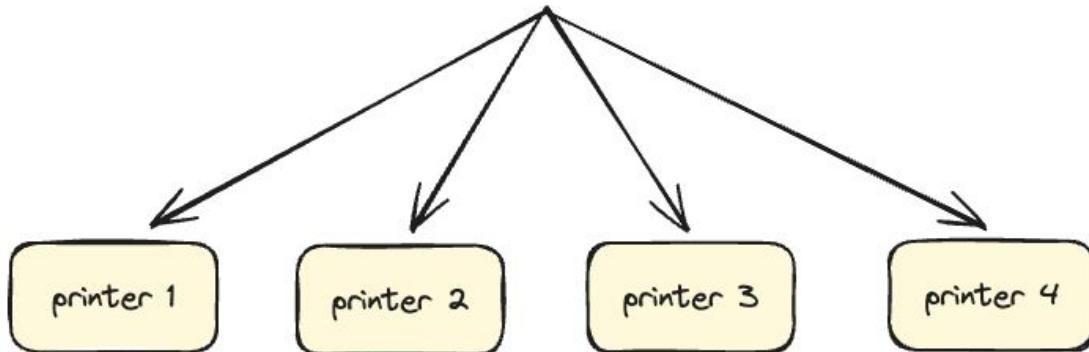


routing

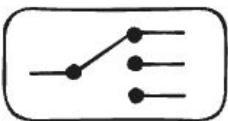


filtering

details	printed
{ename: Orafo, address: Florence ...}	true
{ename: Michael, address: Torino ...}	true
{ename: Gustavo, address: Rome ...}	false
{ename: Ambra, address: Bolzano, ...}	false
...	...
$n > \text{XXL}$	false



race conditions



routing

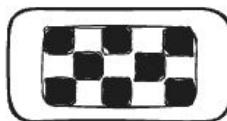
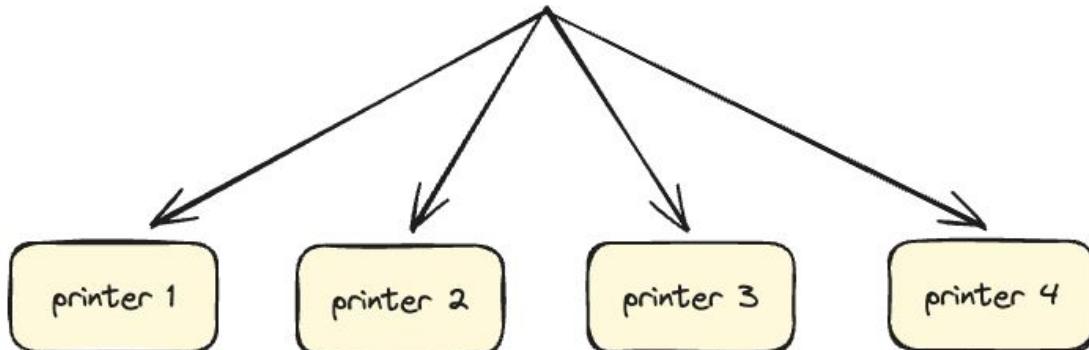


filtering

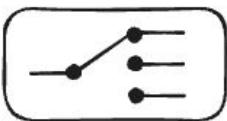


retrying

details	printed
{ename: Orafo, address: Florence ...}	true
{ename: Michael, address: Torino ...}	true
{ename: Gustavo, address: Rome ...}	false
{ename: Ambra, address: Bolzano, ...}	false
...	...
$n > \text{XXL}$	false



race conditions



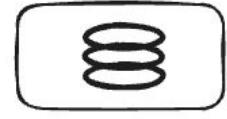
routing



filtering

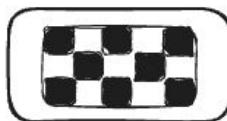
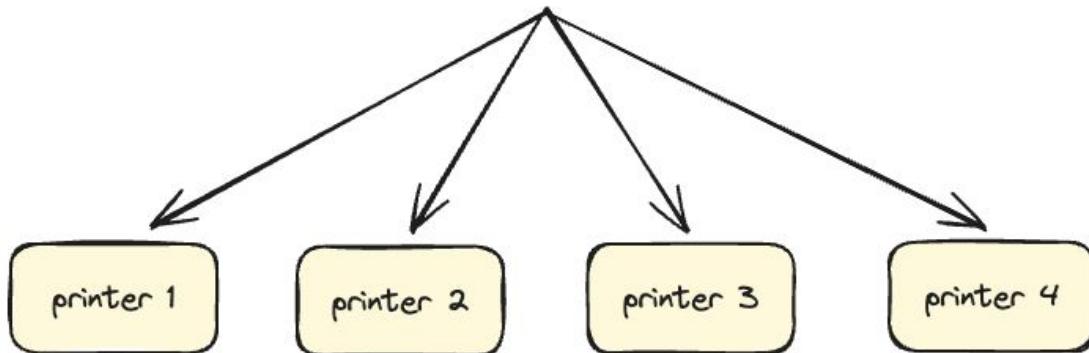


retrying

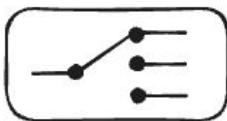


persistence

details	printed
{ename: Orafo, address: Florence ...}	true
{ename: Michael, address: Torino ...}	true
{ename: Gustavo, address: Rome ...}	false
{ename: Ambra, address: Bolzano, ...}	false
...	...
$n > \text{XXL}$	false



race conditions



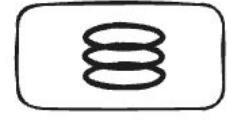
routing



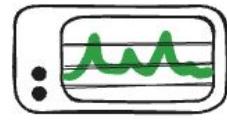
filtering



retrying



persistence



monitoring

# *Message brokers*



Cloud Pub/Sub



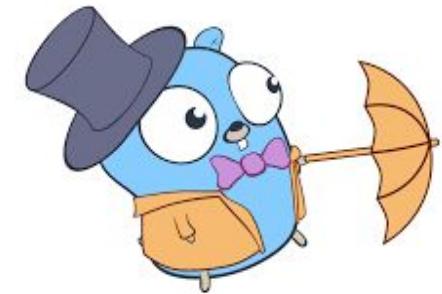
# PGQ message broker



# Broker in the middle

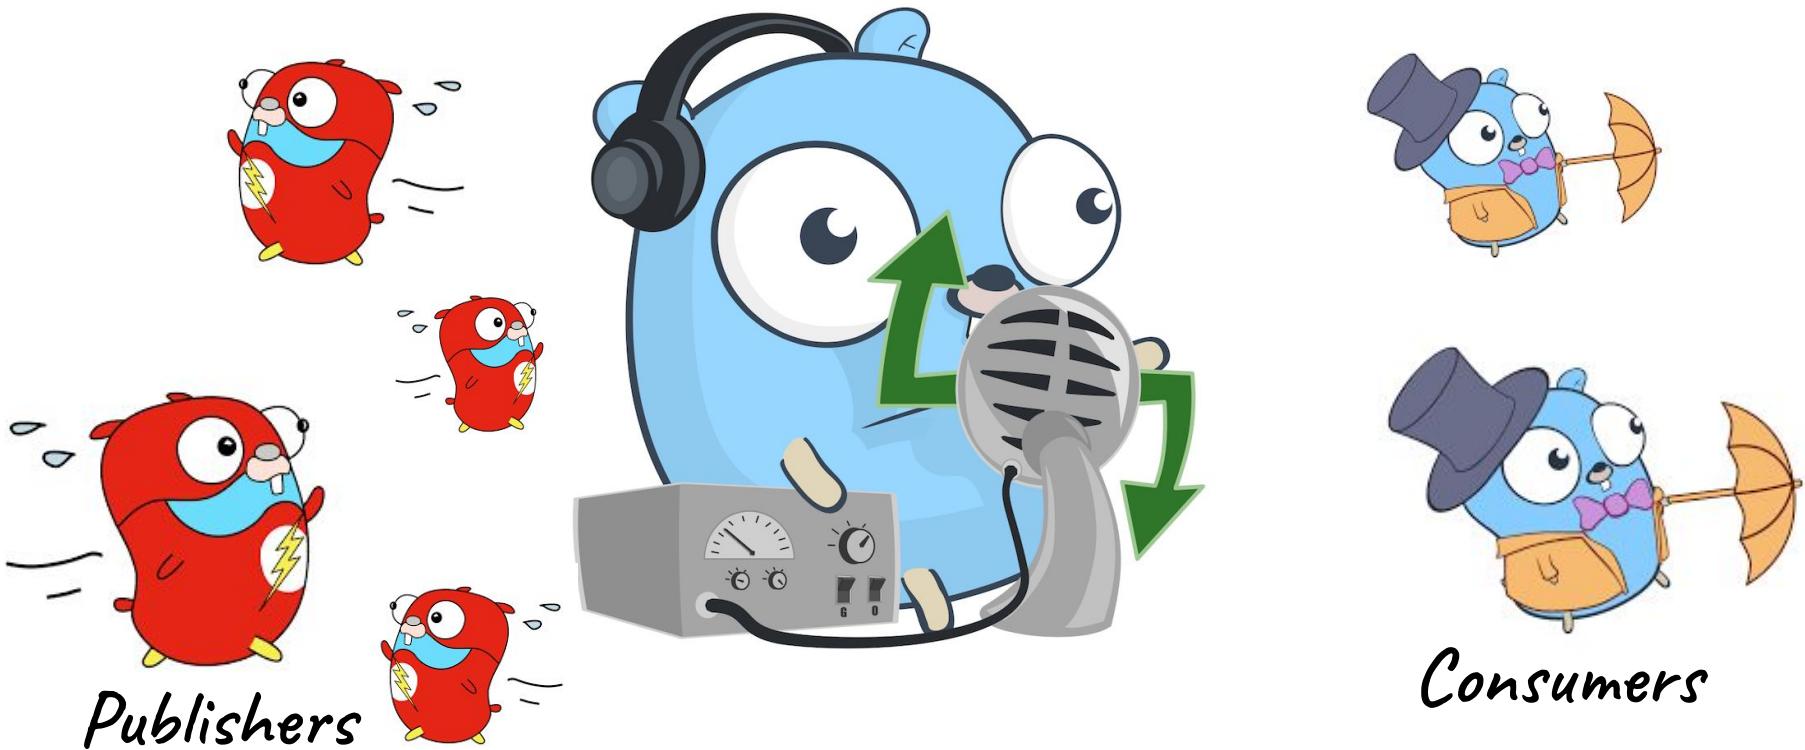


Publisher

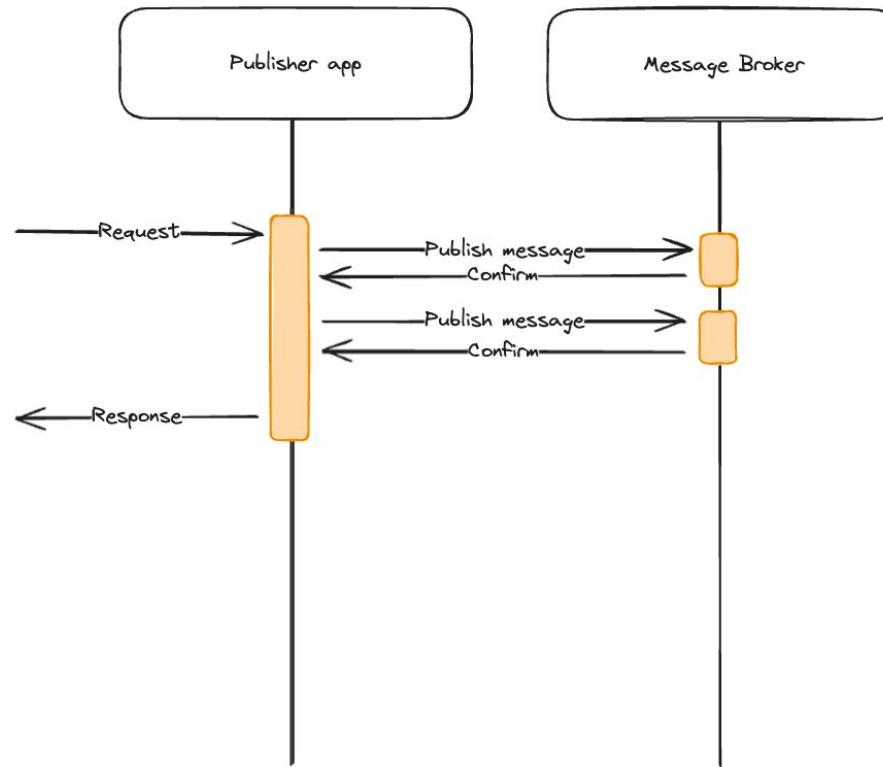


Consumer

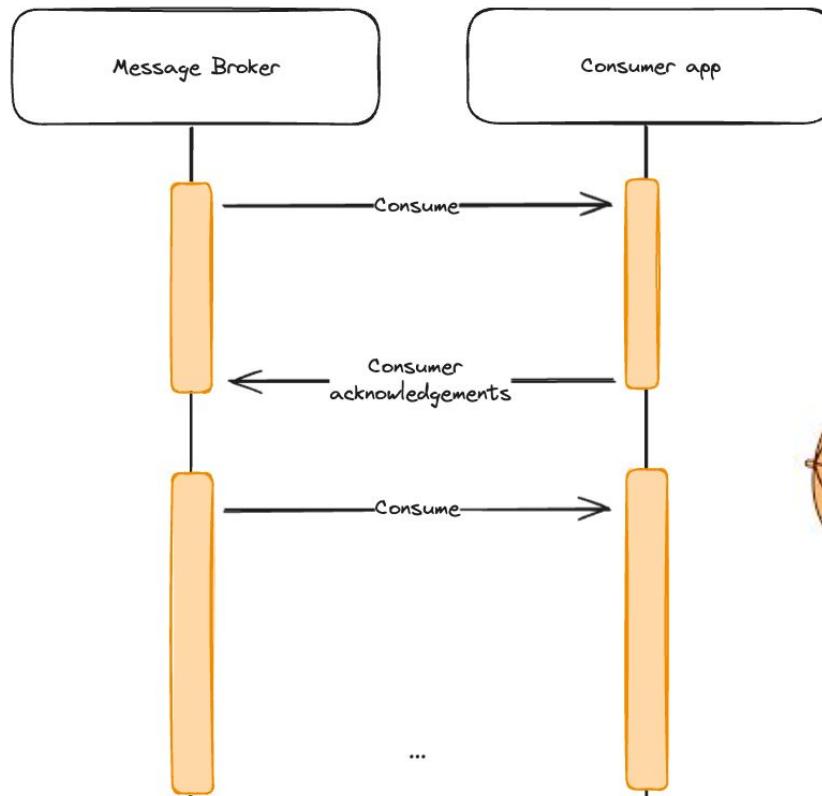
# Broker in the middle



# Publish

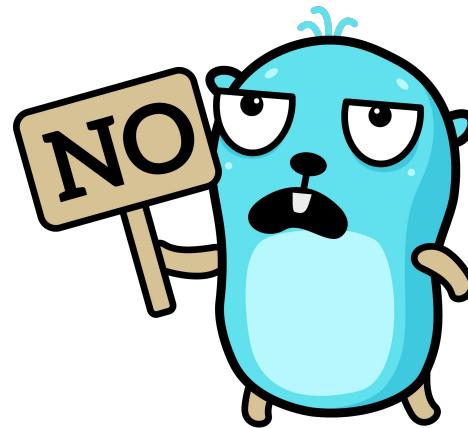
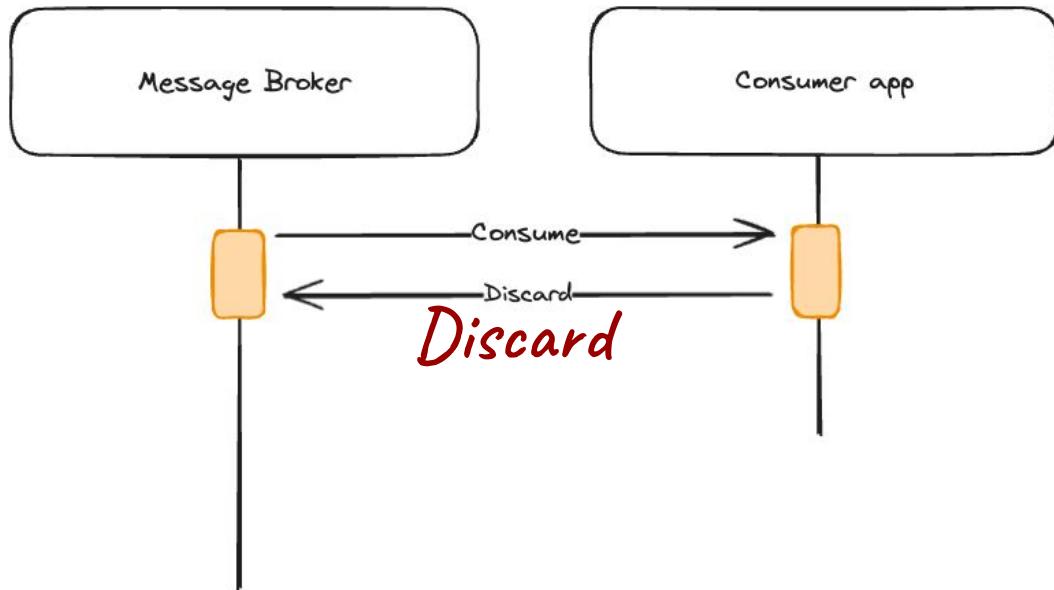


# Consume / Subscribe

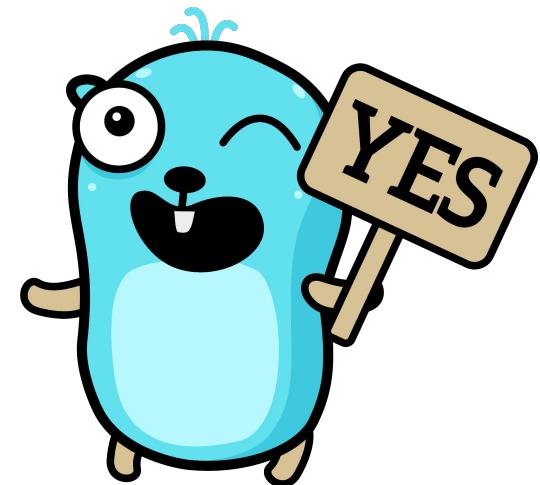
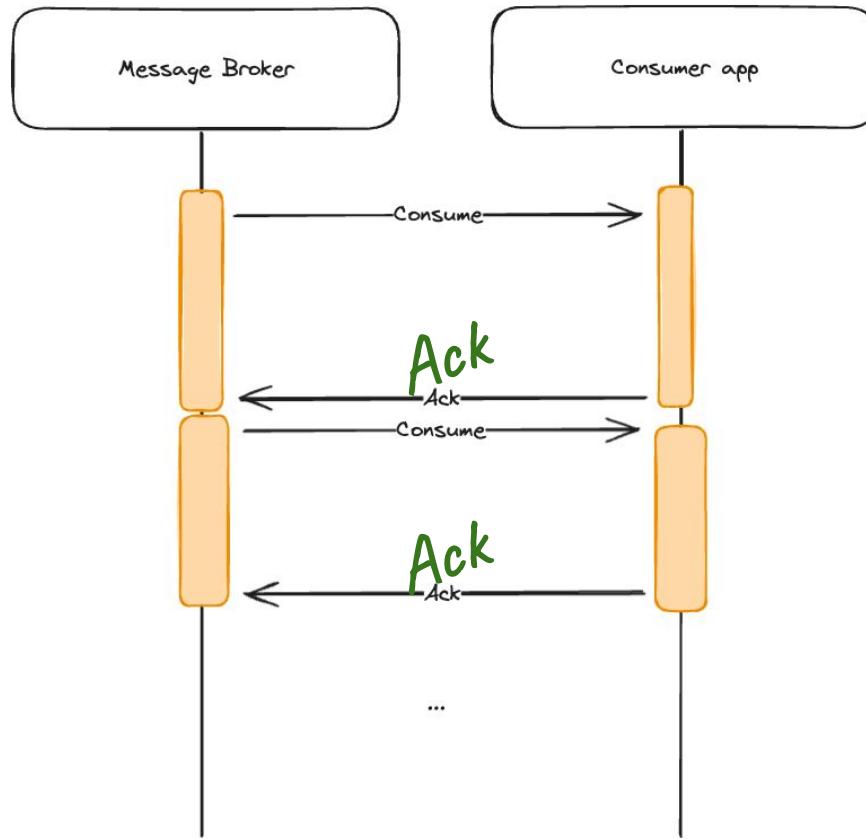


# *Consumer Acknowledgements*

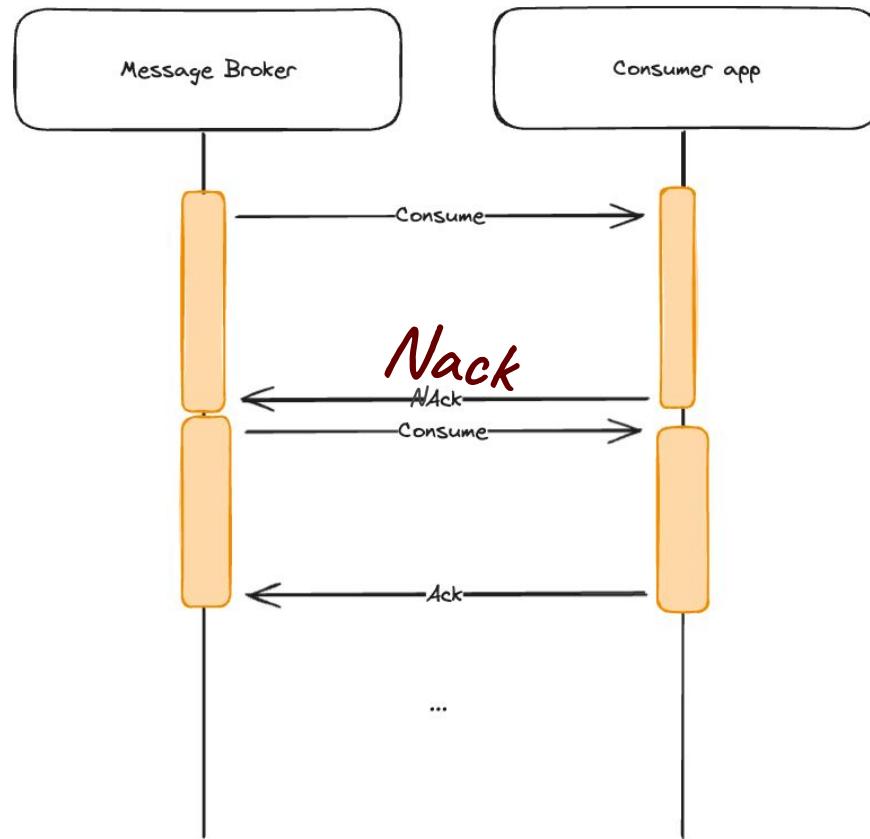
*Discard = will not process the message*  
*Reject*



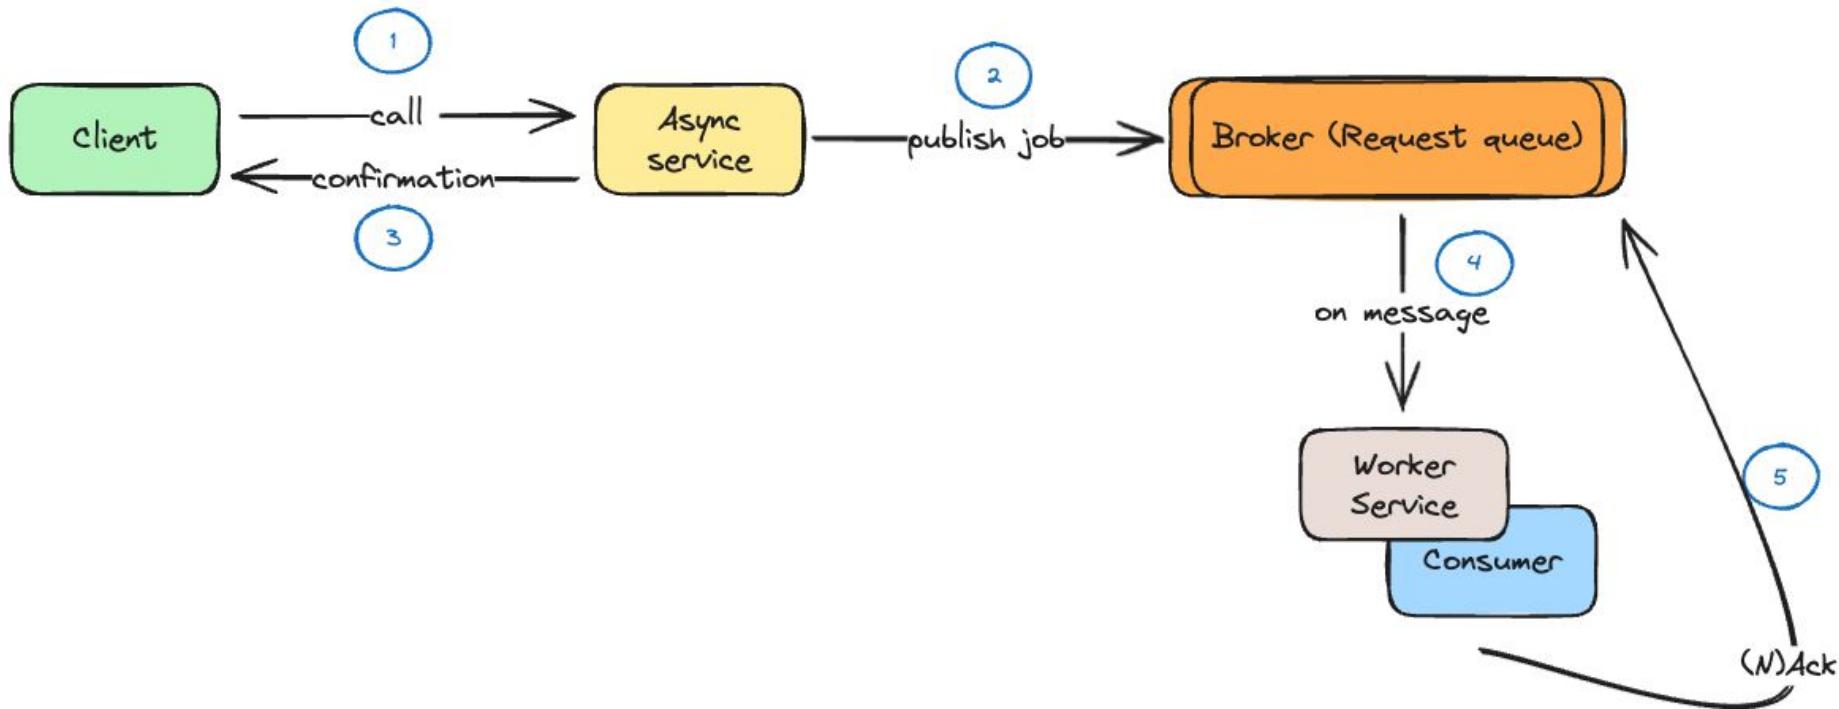
*Acknowledge = I processed the message*



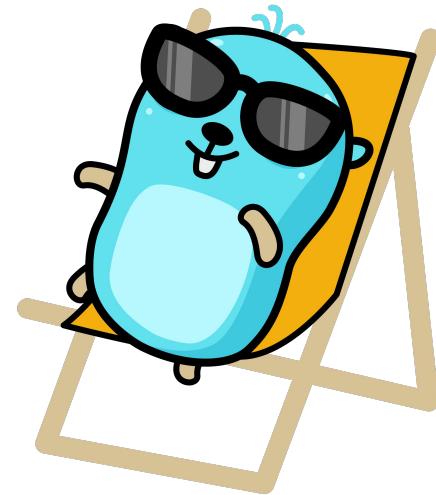
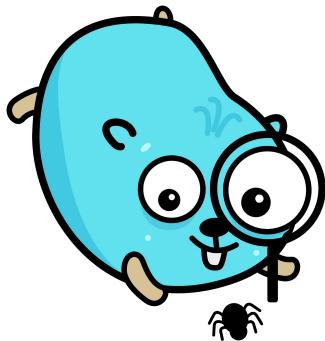
*Negative Acknowledge = I couldn't process the message*



# Request processing

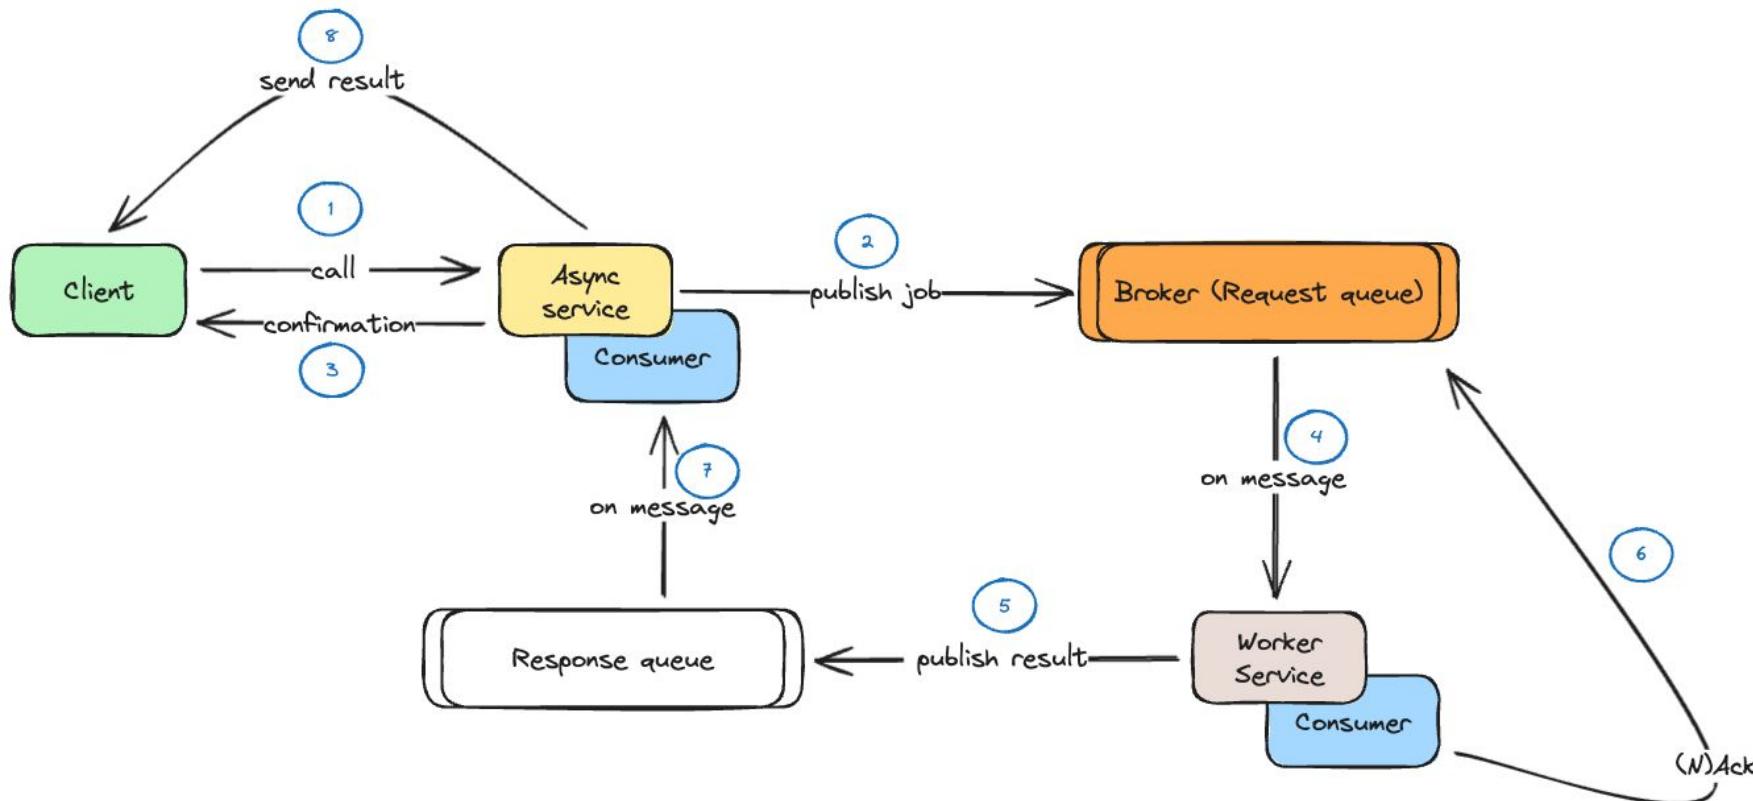


async/rabbit/printer



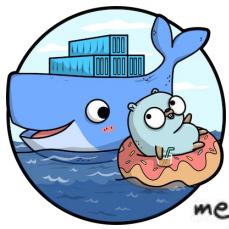
DEMO

# Request-response processing

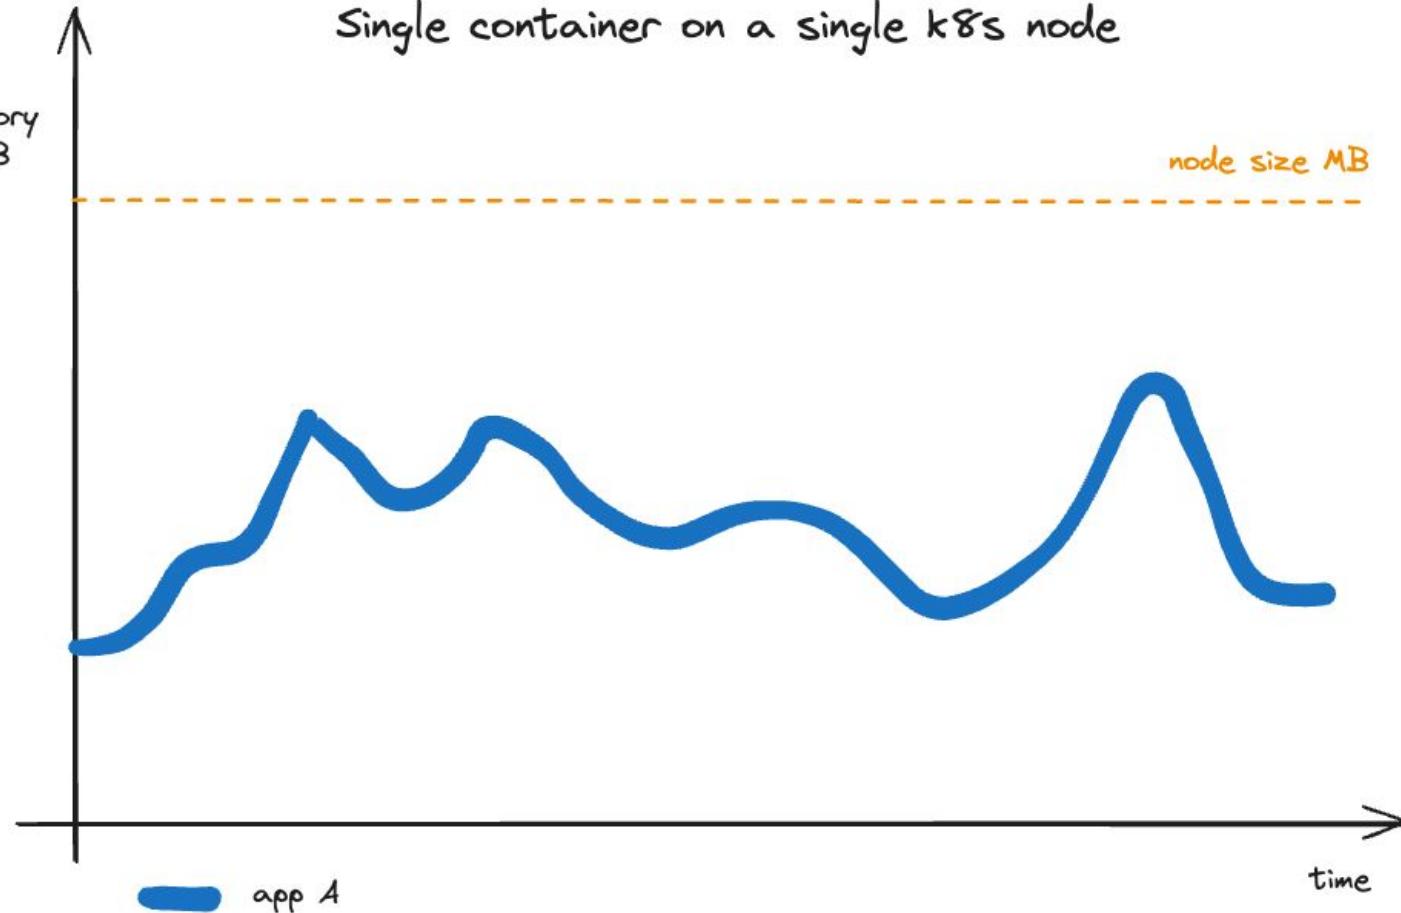


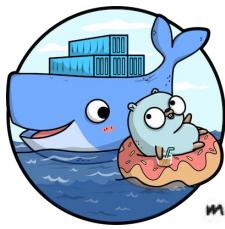
*The theory is nice,  
but the pitfalls reveal the practice*



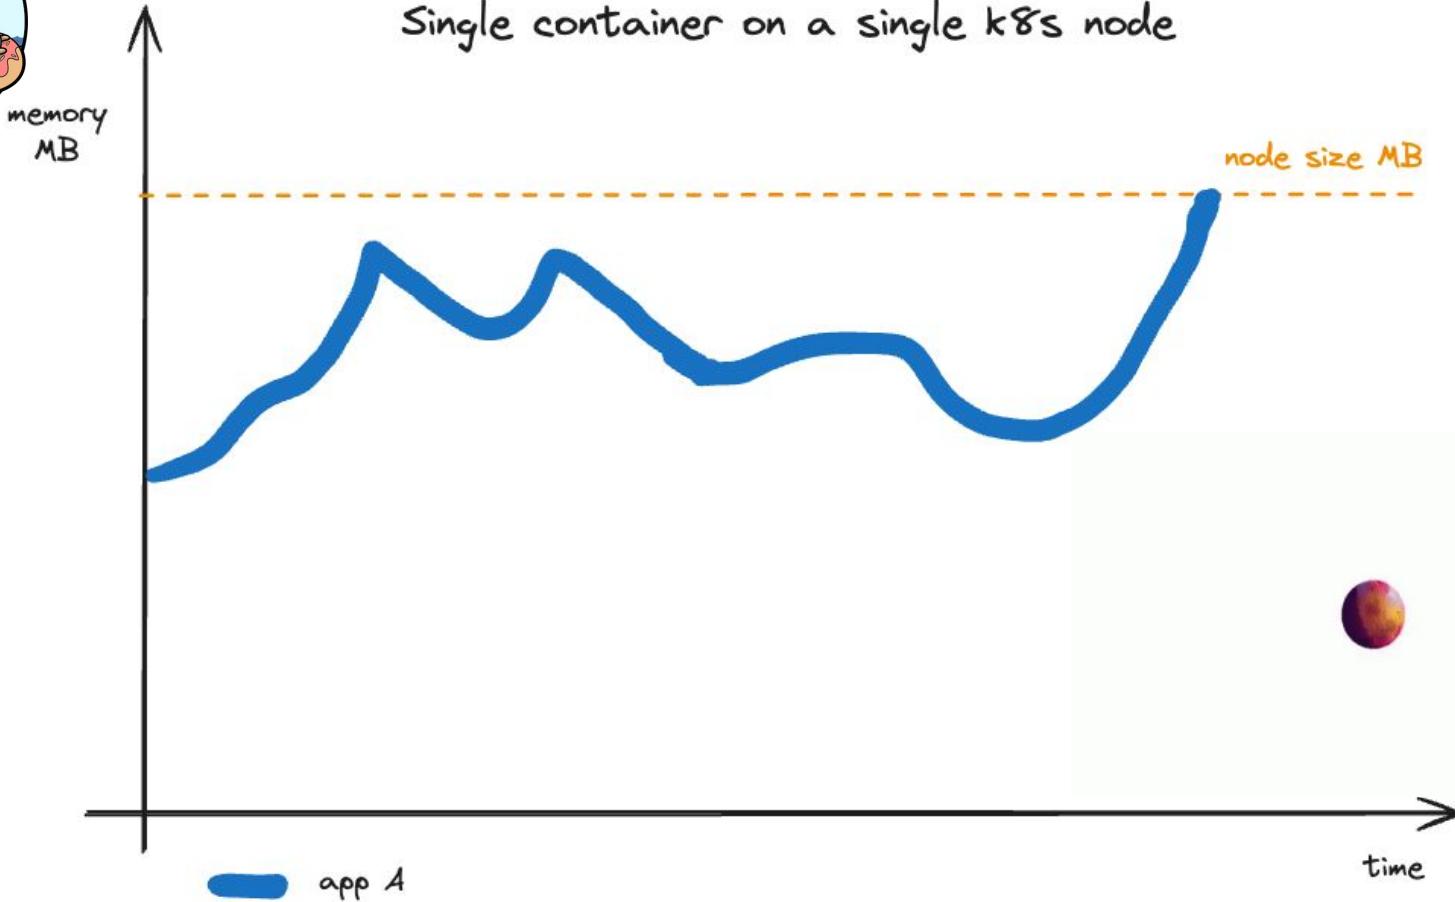


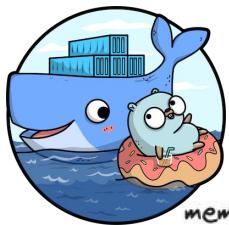
Single container on a single k8s node



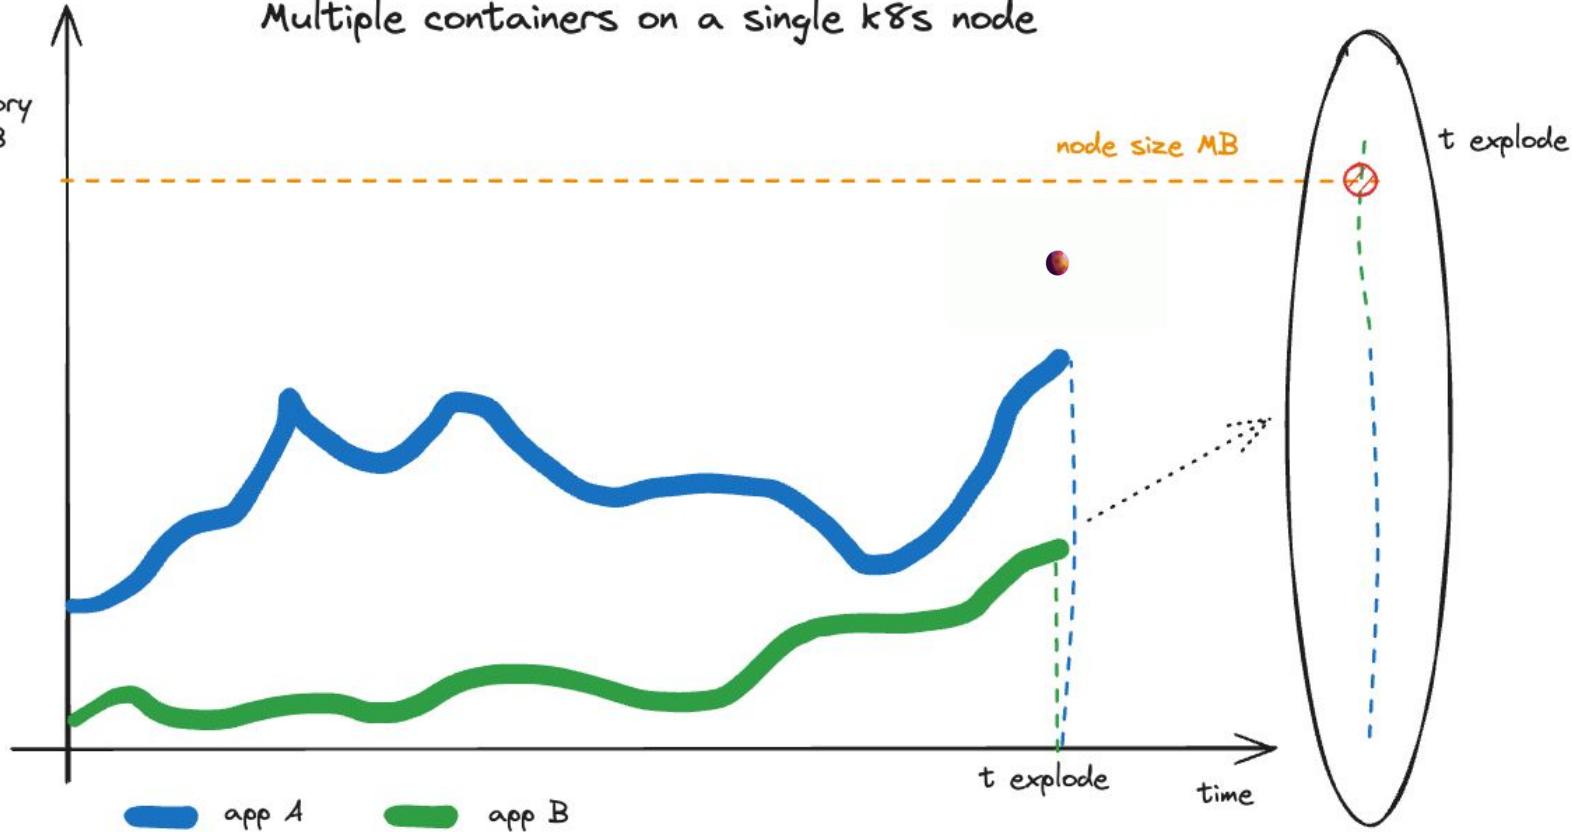


Single container on a single k8s node

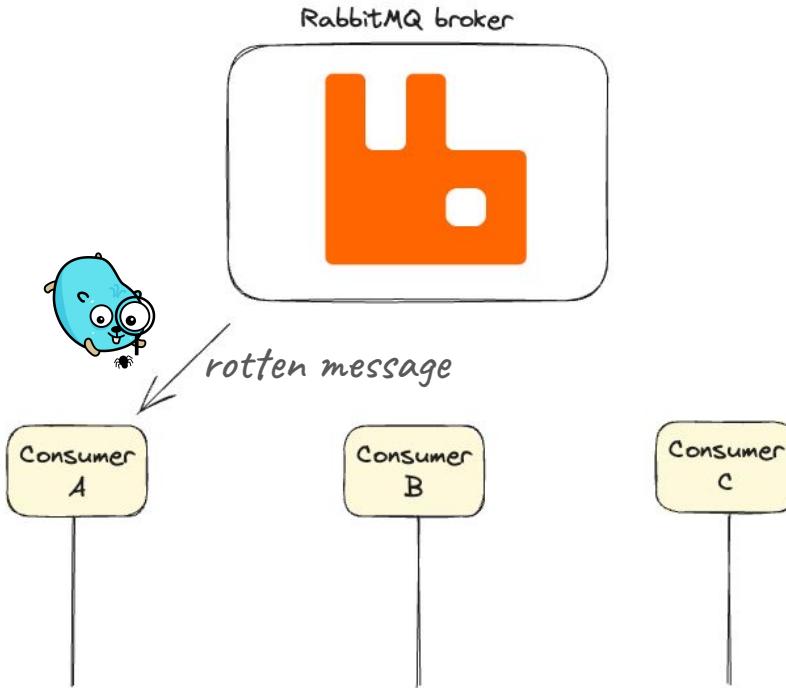




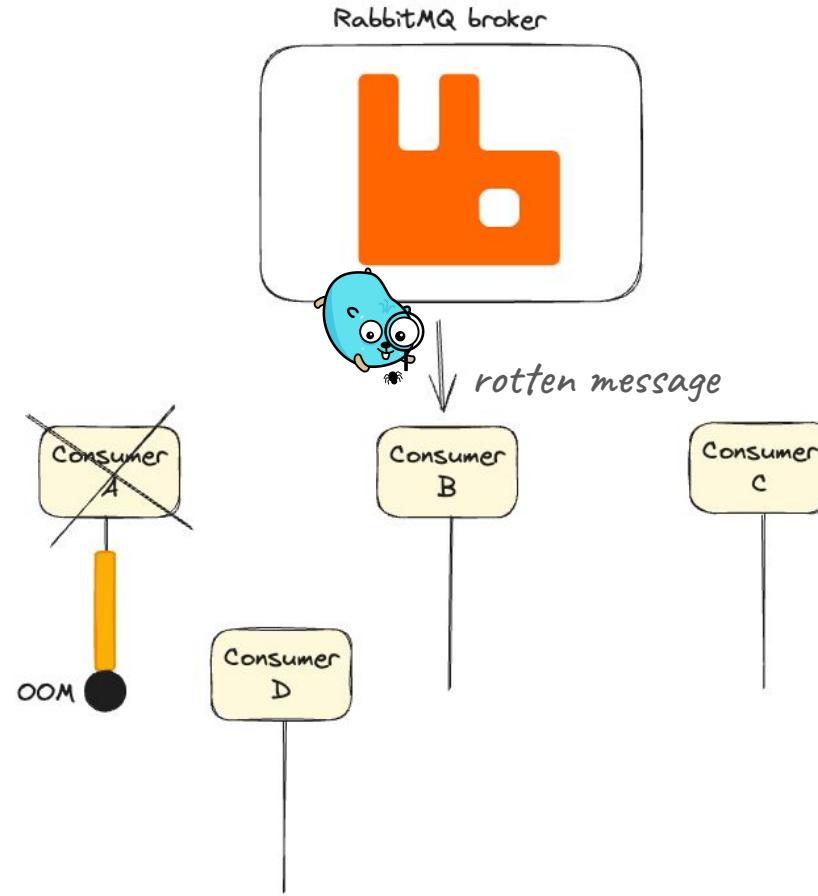
## Multiple containers on a single k8s node



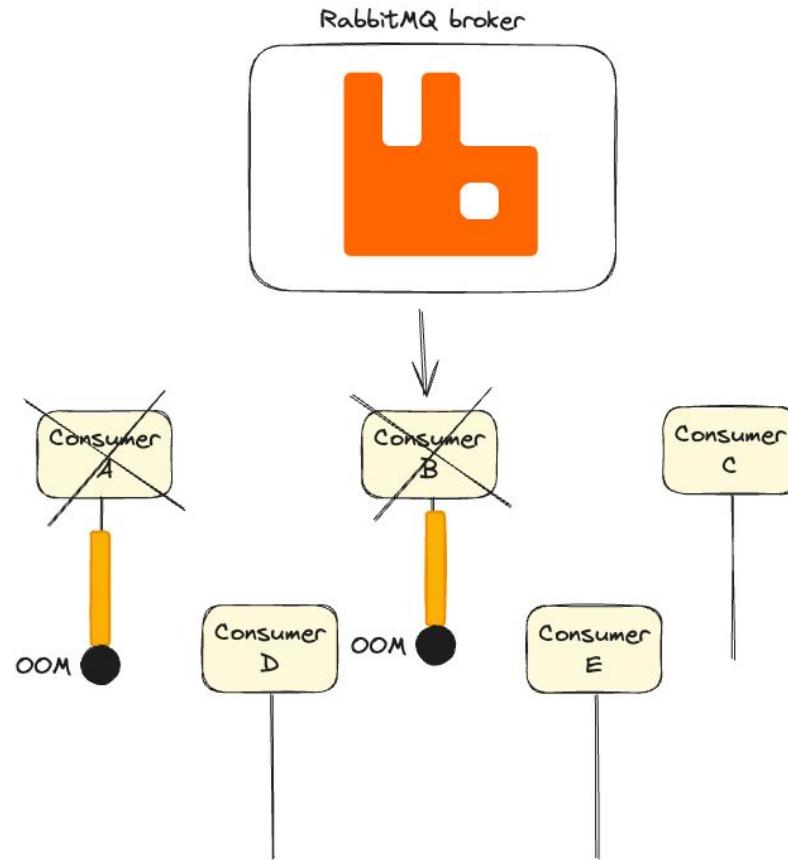
OOM



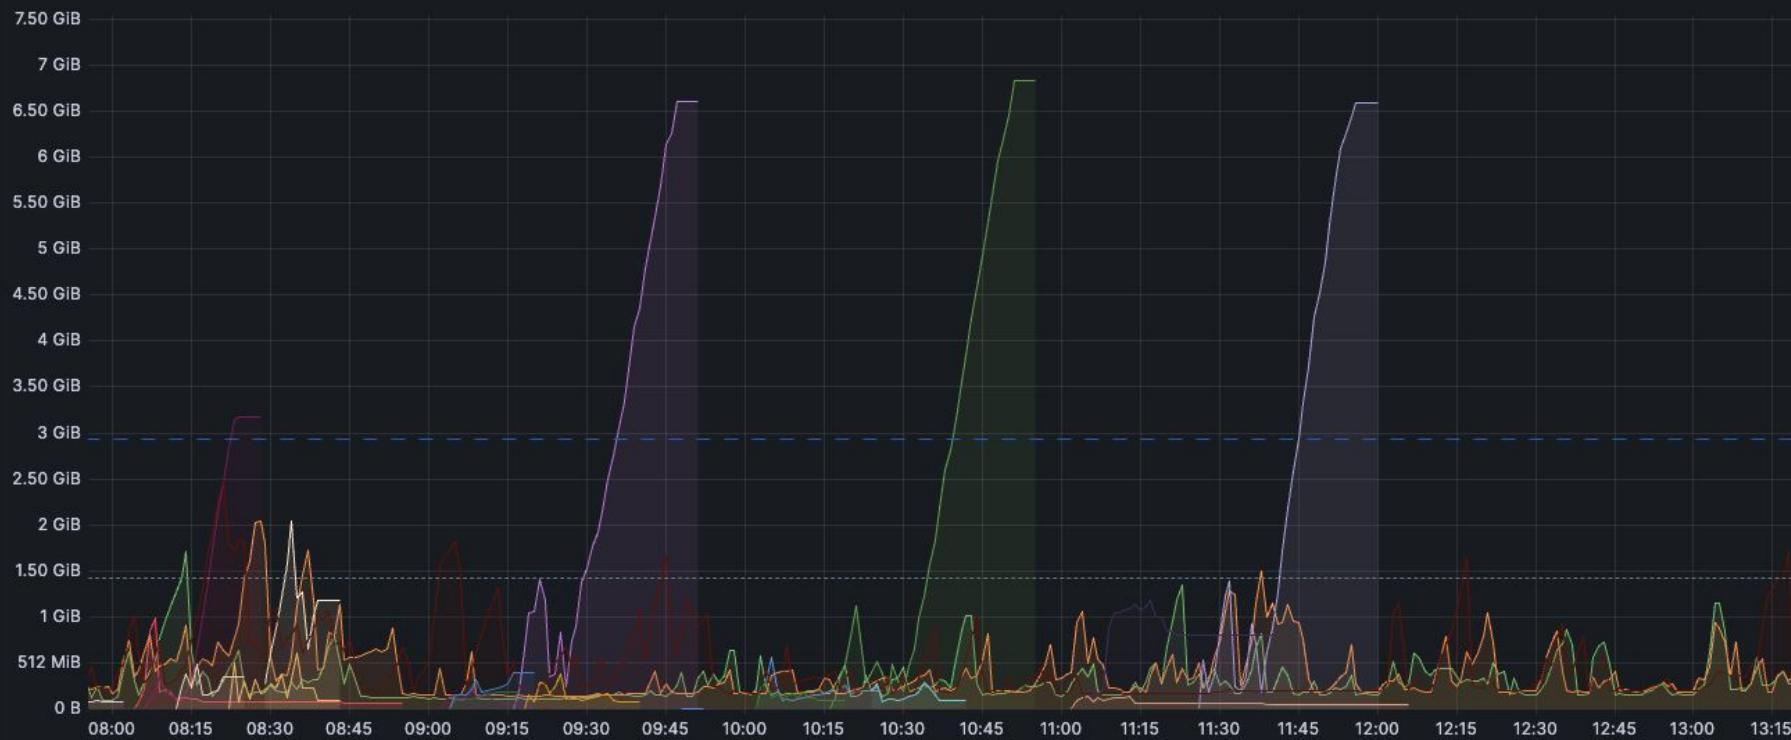
*OOM*



*OOM*

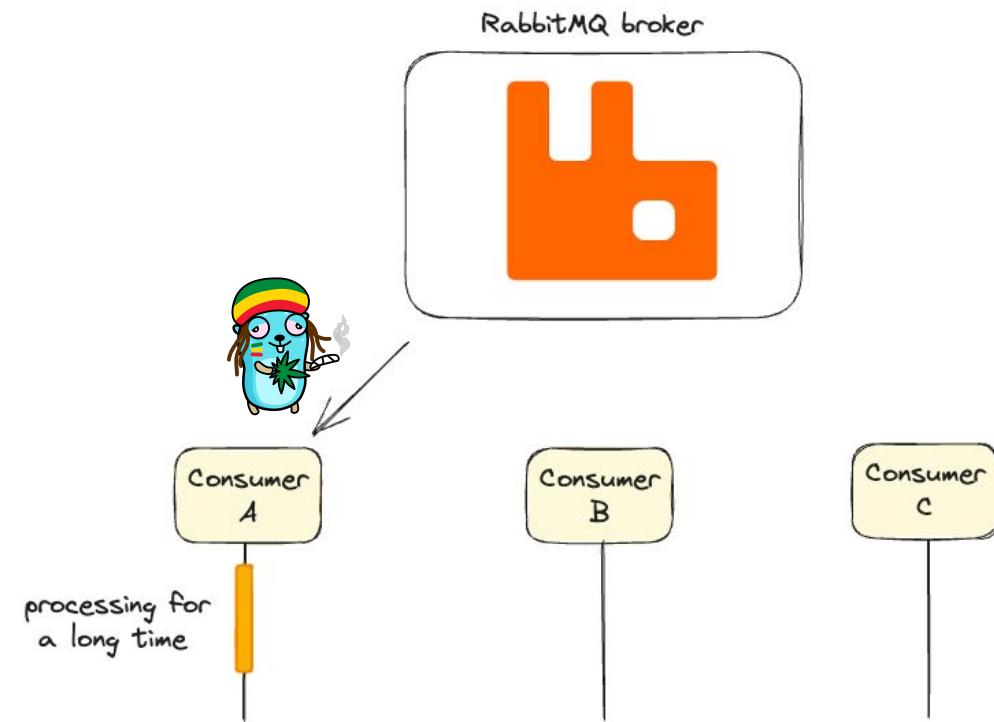


### Extractor mem usage

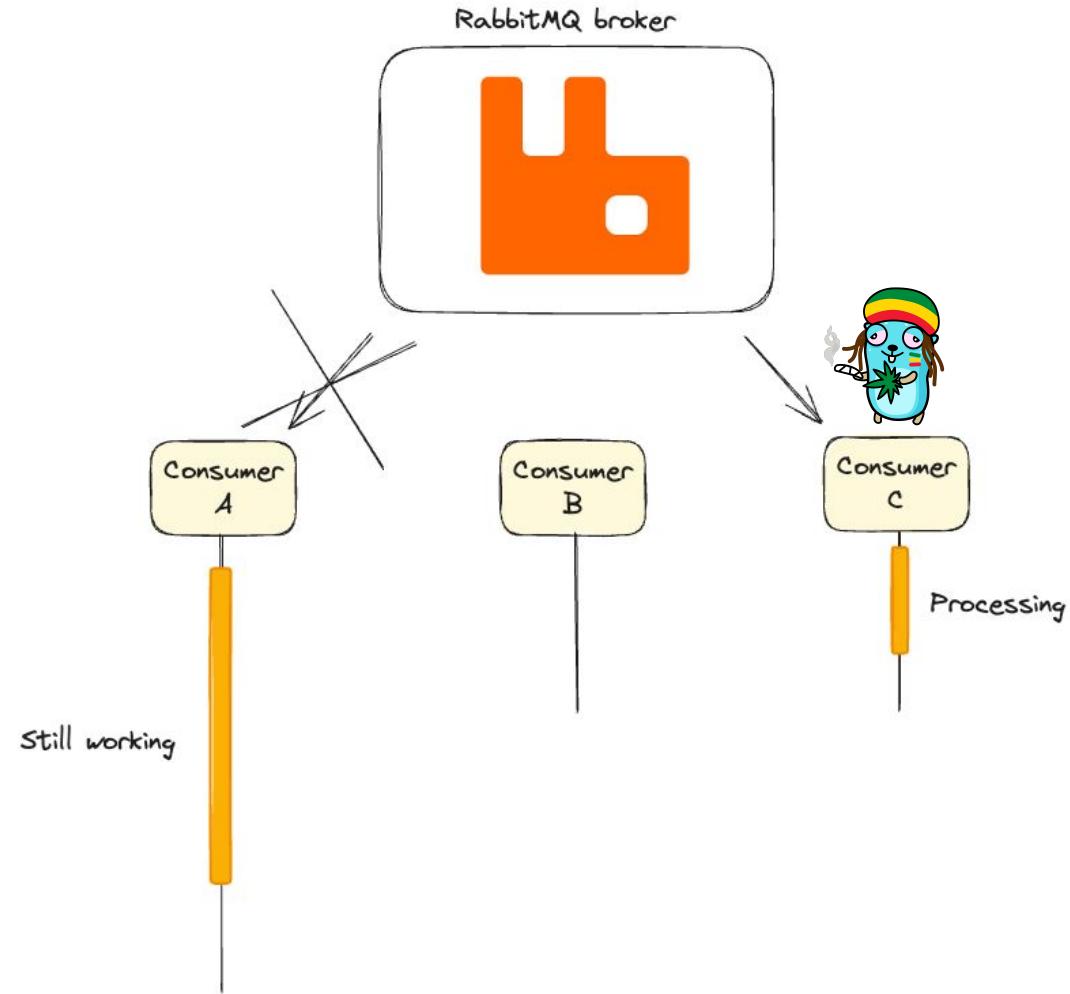


*Deadlock on Delivery Acknowledgement Timeout*

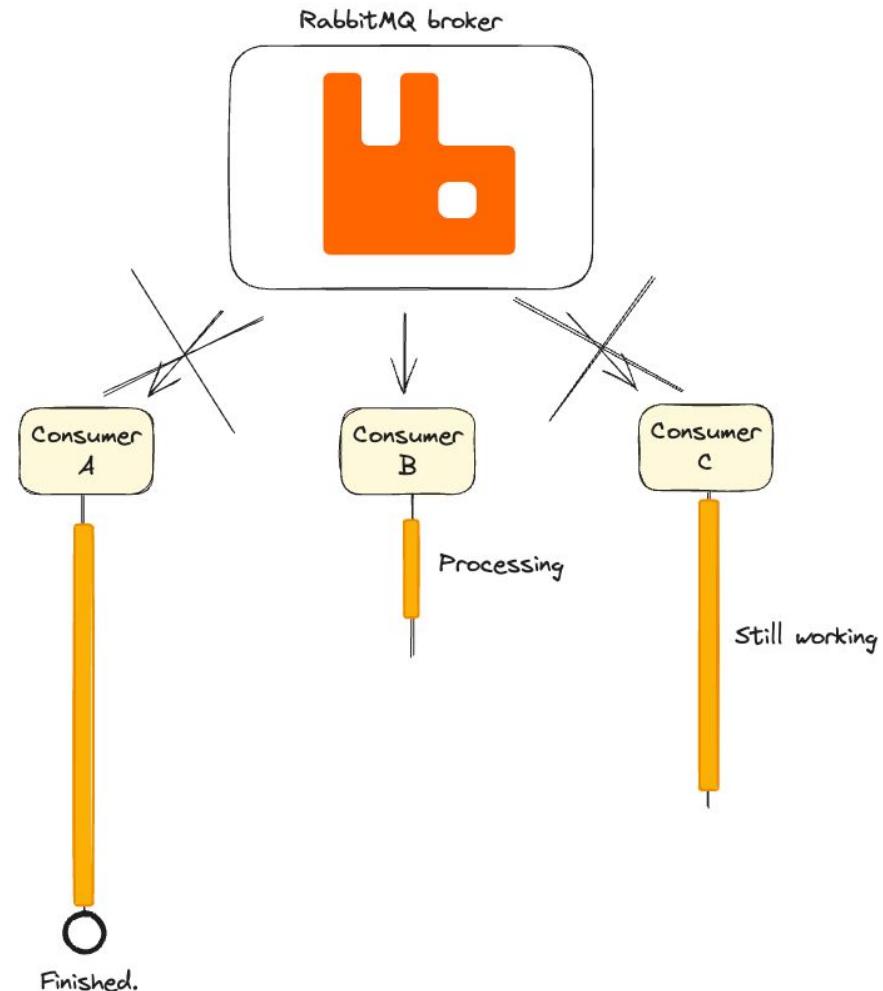
# Ack timeout



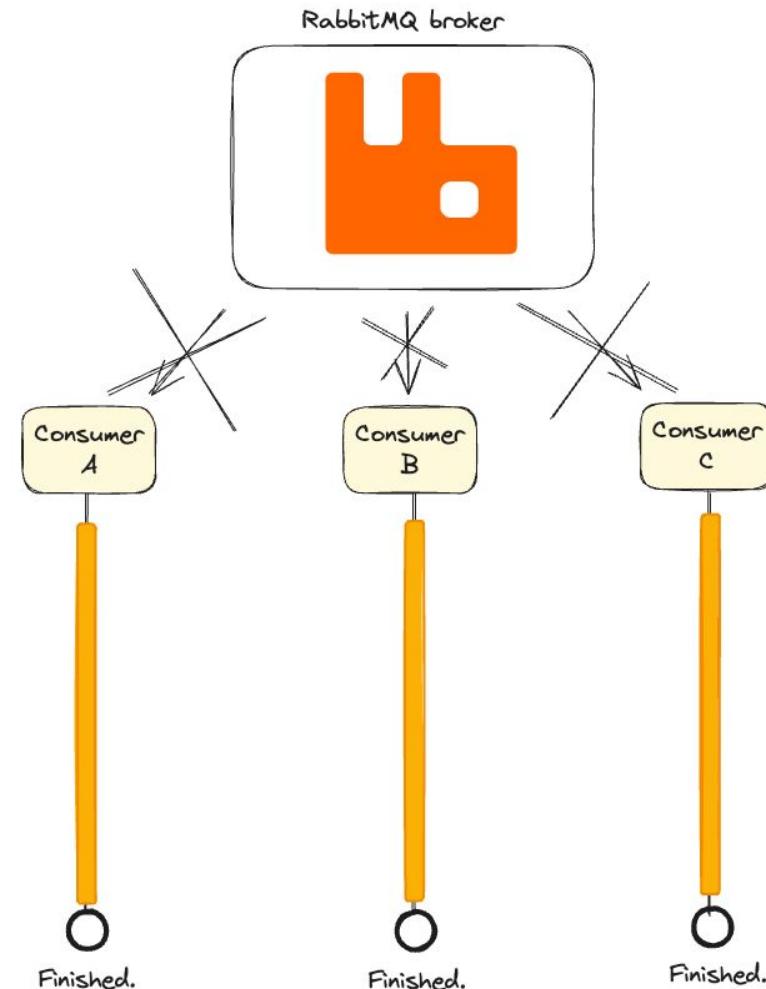
# Ack timeout



# Ack timeout



# Ack timeout





async/pgg



DEMO

*Deciding sync or async*

*It depends.*

Where is the  
asynchronicity helpful?

# Where is the asynchronicity helpful?

## Order Processing Systems

Handling e-commerce orders, where each order might involve multiple steps like payment processing, updating external systems, shipping and notifications, which can be done asynchronously.

# Where is the asynchronicity helpful?

Order Processing Systems

Background Jobs and Task Queues

Tasks such as email sending, file processing, or generating reports, which can be offloaded to a queue and processed by worker services.

# Where is the asynchronicity helpful?

Order Processing Systems

Background Jobs and Task Queues

Decoupling microservices

When microservices need to communicate without waiting for each other, allowing for better decoupling and improved resilience of the system.

# Where is the asynchronicity helpful?

Order Processing Systems

Background Jobs and Task Queues

Decoupling microservices

Data Ingestion and ETL (Extract, Transform, Load) Pipelines

Ingesting large volumes of data from various sources and processing it can be done asynchronously to handle high throughput and avoid blocking.

# Where is the asynchronicity helpful?

Order Processing Systems

Background Jobs and Task Queues

Decoupling microservices

Data Ingestion and ETL (Extract, Transform, Load) Pipelines

Load Balancing and Scaling

Distributing workloads across multiple servers or instances to handle varying load more effectively.

# Where is the asynchronicity helpful?

Order Processing Systems

Background Jobs and Task Queues

Decoupling microservices

Data Ingestion and ETL (Extract, Transform, Load) Pipelines

Load Balancing and Scaling

Distributing workloads across multiple servers or instances to handle varying load more effectively.

Real time notifications, logging, event-driven architectures, ...

# More resources

- Messaging Patterns - Enterprise Integration Patterns
- Watermill website
- Dataddo PGQ package
- Gopher icons

*I got OOM*

*But I am happy to answer your questions.*

*Q&A*









# *Buffer slides*

*Following slides probably will not used at all.*

*Synchronous or Asynchronous?*

# *Deciding sync or async*

## *Blocking vs. Non-blocking*

*If the result is needed immediately to proceed, a synchronous call makes sense*

# Deciding sync or async

## Blocking vs. Non-blocking

If the result is needed immediately to proceed, a synchronous call makes sense

## Responsiveness

In applications with user interfaces, synchronous jobs block the main thread.

# Deciding sync or async

## Blocking vs. Non-blocking

If the result is needed immediately to proceed, a synchronous call makes sense

## Responsiveness

In applications with user interfaces, synchronous jobs block the main thread.

## Control flow

Synchronous calls ensure the tasks are executed in the order.

Asynch. calls can reduce the execution time, but their coordination may be tricky.

# Deciding sync or async

## Blocking vs. Non-blocking

If the result is needed immediately to proceed, a synchronous call makes sense

## Responsiveness

In applications with user interfaces, synchronous jobs block the main thread.

## Control flow

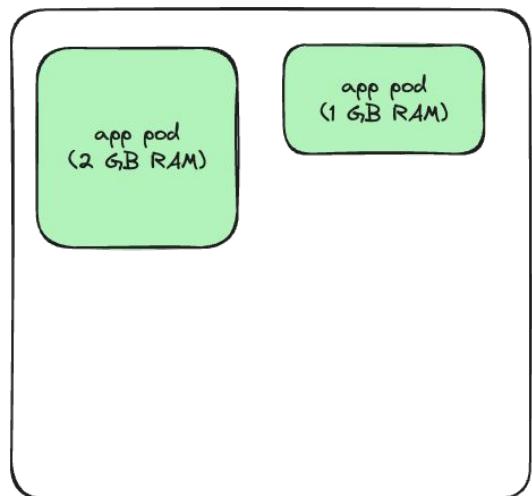
Synchronous calls ensure the tasks are executed in the order.

Asynch. calls can reduce the execution time, but their coordination may be tricky.

## Scalability

Asynchronous processing can help scale more effectively (free resources)

k8s node (8 GiB RAM)



k8s node (8 GiB RAM)

k8s node (8 GiB RAM)

app pod  
(2 GiB RAM)

app pod  
(1 GiB RAM)

app pod  
(4 GiB RAM)

app pod  
(3 GiB RAM)

# OOM

